

CMPE 411

Computer Architecture

Lecture 28

Introduction to Multiprocessor Systems

December 7, 2017

www.csee.umbc.edu/~younis/CMPE411/CMPE411.htm



Lecture's Overview

□ Previous Lecture

- ➔ Interfacing I/O devices
 - Interfacing with I/O devices
 - Communication with I/O devices
 - Operating System's role
- ➔ Memory to processor interconnect
 - Definition of bus structure
 - Bus transactions
 - Types of buses
 - Bus Standards
- ➔ Bus Performance and Protocol
 - Synchronous versus Asynchronous buses
 - Bandwidth optimization factors
 - Single versus multiple master bus
 - Bus arbitration approaches

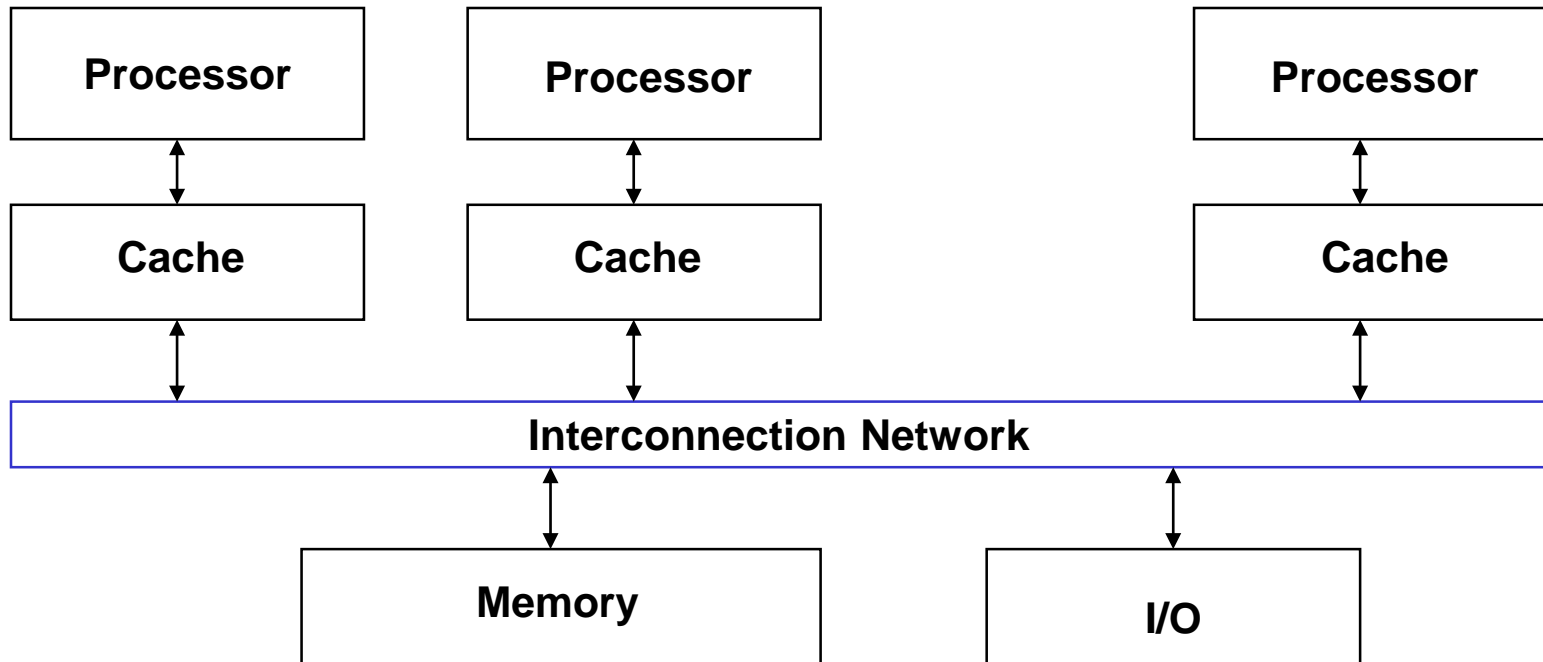
□ This Lecture

- ➔ Introduction to multiprocessor systems
- ➔ Issues of parallelism
- ➔ Architecture models for multiprocessors



The Big Picture: Where are We Now?

❑ **Multiprocessor** – a computer system with at least two processors



- ➔ Can deliver high throughput for independent jobs via **job-level parallelism** or **process-level parallelism**
- ➔ And improve the run time of a *single* program that has been specially crafted to run on a multiprocessor - a **parallel processing program**

Popular Multiprocessor Systems

Multicores Now Common

- ❑ The power challenge has forced a change in the design of microprocessors
 - ➔ Since 2002 the rate of improvement in the response time of programs has slowed from a factor of 1.5 per year to less than a factor of 1.2 per year
- ❑ Today's microprocessors typically contain more than one core – **Chip Multicore microProcessors (CMPs)** – in a single IC
 - ➔ The number of cores is expected to double every two years

Product	AMD Barcelona	Intel Nehalem	IBM Power 6	Sun Niagara 2
Cores per chip	4	4	2	8
Clock rate	2.5 GHz	~2.5 GHz?	4.7 GHz	1.4 GHz
Power	120 W	~100 W?	~100 W?	94 W

Computing Cluster

- ❑ A **cluster** – a set of independent servers (or PCs) connected over a local area network (LAN) functioning as a single large multiprocessor
 - ➔ Search engines, Web servers, email servers, databases, ...
- ❑ A key challenge is to craft parallel (concurrent) programs that have high performance on multiprocessors as the number of processors increase:
 - ➔ Load balancing, time for synchronization, overhead for communication

* Slide is a courtesy of Mary Jane Irwin




Encountering Amdahl's Law

□ Speedup due to enhancement E is

$$\text{Speedup w/ E} = \frac{\text{Exec time w/o E}}{\text{Exec time w/ E}}$$

- Suppose that enhancement E accelerates a fraction F (F < 1) of the task by a factor S (S > 1) and the remainder of the task is unaffected



$$\text{ExTime w/ E} = \text{ExTime w/o E} \times ((1-F) + F/S)$$

$$\text{Speedup w/ E} = 1 / ((1-F) + F/S)$$

Example 1: Amdahl's Law

$$\text{Speedup w/ } E = 1 / ((1-F) + F/S)$$

- ❑ Consider an enhancement which runs 20 times faster but which is only usable 25% of the time.

$$\text{Speedup w/ } E = 1 / (.75 + .25/20) = 1.31$$

- ❑ What if its usable only 15% of the time?

$$\text{Speedup w/ } E = 1 / (.85 + .15/20) = 1.17$$

- ❑ Amdahl's Law tells us that to achieve linear speedup with 100 processors, **none** of the original computation can be scalar!
- ❑ To get a speedup of 90 from 100 processors, the percentage of the original program that could be scalar would have to be 0.1% or less

$$\text{Speedup w/ } E = 1 / (.001 + .999/100) = 90.99$$

Opportunities for Applications

Scientific Computing

- ❑ Nearly Unlimited Demand (Grand Challenge):

<i>App</i>	<i>Perf (GFLOPS)</i>	<i>Memory (GB)</i>
48 hour weather	0.1	0.1
72 hour weather	3	1
Pharmaceutical design	100	10
Global Change, Genome	1000	1000

- ❑ Successes in some real industries:

- ➔ Petroleum: reservoir modeling
- ➔ Automotive: crash simulation, drag analysis, engine
- ➔ Aeronautics: airflow analysis, engine, structural mechanics
- ➔ Pharmaceuticals: molecular modeling
- ➔ Entertainment: full length movies (“Toy Story”)

Commercial Applications

- ❑ Transaction processing, File servers, electronic CAD simulation (multiple processes), WWW search engines
- ❑ Examples: IBM RS6000, Tandem (Compaq) Himalaya



* Slide is a courtesy of Dave Patterson

Parallel Computers

- ❑ Definition: “A parallel computer is a collection of processing elements that cooperate and communicate to solve large problems fast.”

Almasi and Gottlieb, *Highly Parallel Computing* ,1989

- ❑ Parallel machines is expected to have a bigger role in the future since:
 - Microprocessors are likely to remain dominant in the uniprocessor arena and the logical way to extend the performance is by connecting multiple microprocessors
 - It is not expected that the microprocessor technology will keep the pace of performance improvement given the increased level of complexity
 - There has been steady progress in software development for parallel architectures in recent years
- ❑ Questions about parallel computers:
 - ➔ How large a collection?
 - ➔ How powerful are processing elements?
 - ➔ How do they cooperate and communicate?
 - ➔ How are data transmitted?
 - ➔ What type of interconnection?
 - ➔ What are HW and SW primitives for programmer?
 - ➔ Does it translate into performance?



* Slide is a courtesy of Dave Patterson

Taxonomy of Parallel Architecture

Flynn Categories

- ❑ SISD (Single Instruction Single Data): Uniprocessors
- ❑ MISD (Multiple Instruction Single Data): No commercial machine yet
- ❑ SIMD (Single Instruction Multiple Data): Vector machine
 - ➔ Examples: Illiac-IV, CM-2
 - Simple programming model, Low overhead, Flexibility
- ❑ MIMD (Multiple Instruction Multiple Data)
 - ➔ Two classes: centralized shared memory and distributed memory architectures
 - ➔ Examples: Sun Enterprise 5000, Cray T3D, SGI Origin
 - Flexible, *Use off-the-shelf micros*

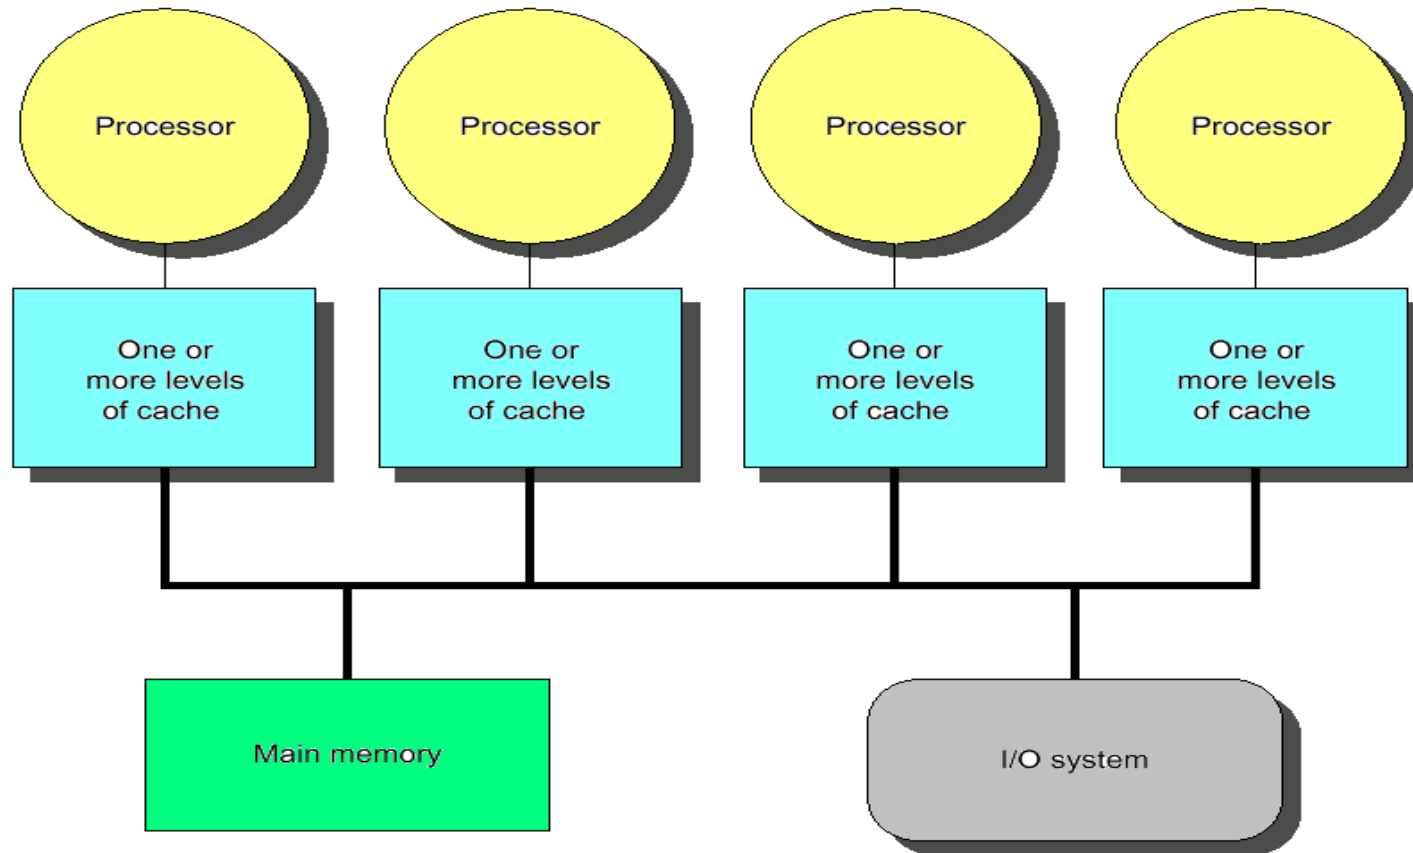
Level of Parallelism

- ❑ Bit level parallelism: 1970 to 1985
 - ➔ 4 bits, 8 bit, 16 bit, 32 bit microprocessors
- ❑ Instruction level parallelism (ILP): 1985 through today
 - ➔ Pipelining, Superscalar, VLIW, Out-of-Order execution
- ❑ Process/Thread level parallelism: mainstream for general purpose computing
 - ➔ Servers are parallel
 - ➔ High-end Desktop multi (dual, quad, ...) processor PC

* Slide is a courtesy of Dave Patterson

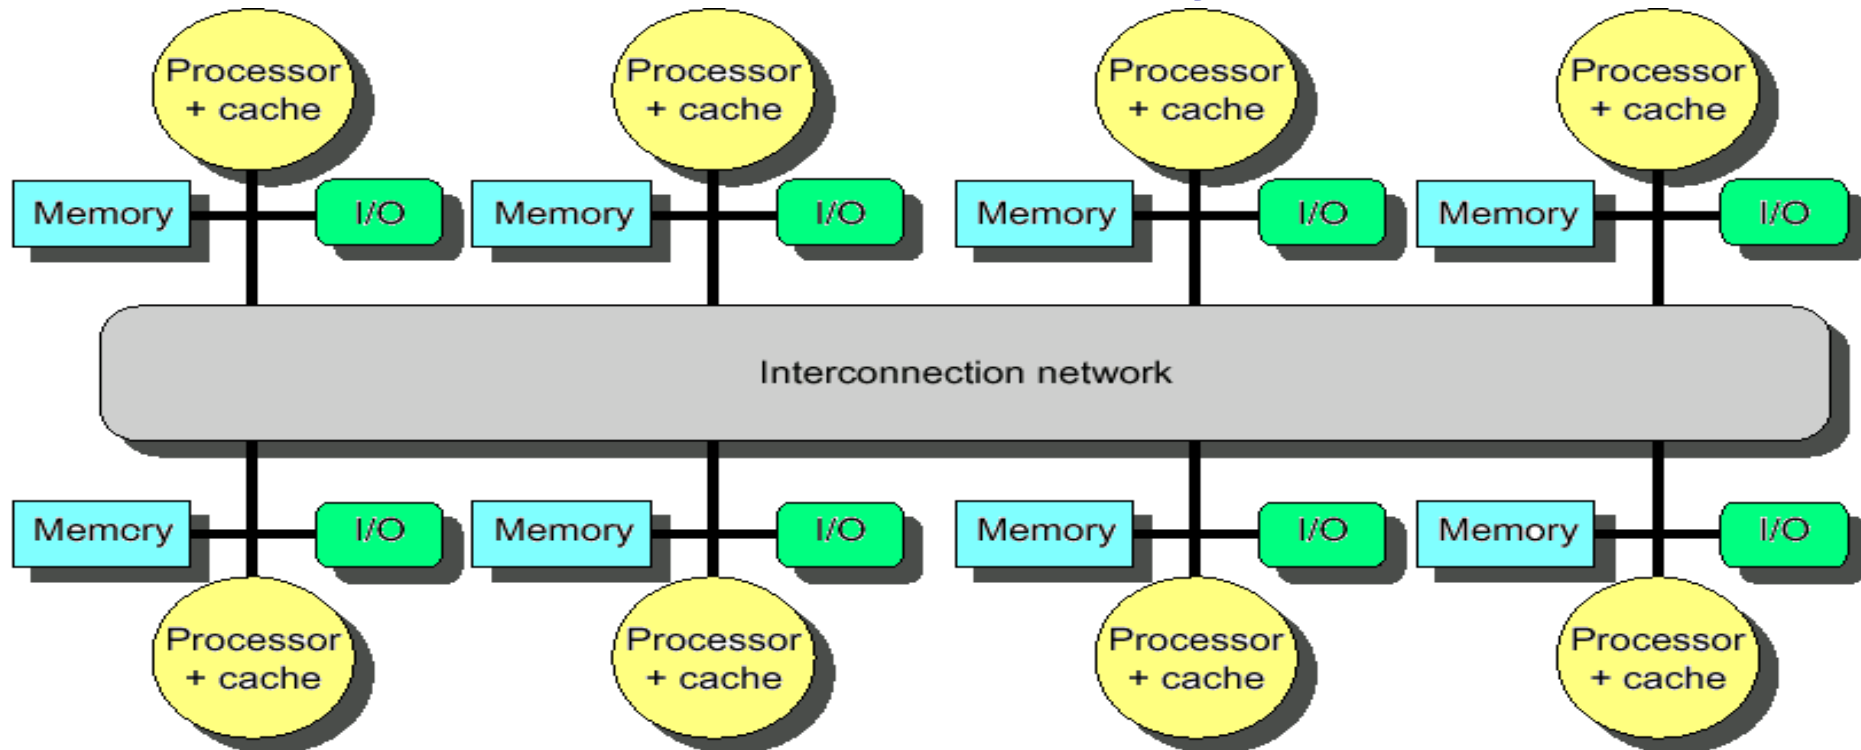


Centralized Shared Memory MIMD



- ❑ Processors share a single centralized memory through a bus interconnect
- ❑ Feasible for small processor count to limit memory contention
- ❑ Usually called “Uniform Memory Access” or UMA model
- ❑ Centralized shared memory architectures are the most common form of MIMD design

Distributed Memory MIMD



- ☐ Uses physically distributed memory to support large processor counts (to avoid memory contention)
- ☐ Processor nodes can have some I/O devices and might be composed of multiple processors in a centralized shared memory architecture (clustering)
- ☐ Distributed memory MIMD provides the following advantages
 - Allows cost-effective way to scale the memory bandwidth
 - Reduces memory latency
- ☐ Increased complexity of communicating data is the major disadvantage

Communication Models

☐ Shared Memory

➔ Processors communicate with shared address space

➔ Easy on small-scale machines

➔ Properties:

- Model of choice for uniprocessors, small-scale multiprocessor
- Ease of programming
- Lower latency
- Easier to use hardware controlled caching
- Difficult to handle node failure

☐ Message passing

➔ Processors have private memories, communicate via messages

➔ Properties:

- Less hardware, easier to design
- Focuses attention on costly non-local operations

☐ Can support either SW models on any HW bases

Summing 100,000 Numbers on 100 Proc.

- ❑ Processors start by running a loop that sums their subset of vector A numbers (vectors A and sum are shared variables, P_n is the processor's number, i is a private variable)

```
sum[Pn] = 0;  
for (i = 1000*Pn; i < 1000*(Pn+1); i = i + 1)  
    sum[Pn] = sum[Pn] + A[i];
```

- ❑ The processors then coordinate in adding together the partial sums (half is a private variable initialized to 100 (the number of processors)) – reduction


```
repeat  
    synch(); /*synchronize first  
    if (half%2 != 0 && Pn == 0)  
        sum[0] = sum[0] + sum[half-1];  
    half = half/2  
    if (Pn<half) sum[Pn] = sum[Pn] + sum[Pn+half]  
until (half == 1); /*final sum in sum[0]
```

* Slide is a courtesy of Mary Jane Irwin



An Example with 10 Processors (Shared Memory Model)

sum[P0]sum[P1]sum[P2] sum[P3]sum[P4]sum[P5]sum[P6] sum[P7]sum[P8] sum[P9]



Summing 100,000 Numbers on 100 Proc. (Message Passing Model)

- Start by distributing 1000 elements of vector A to each of the local memories and summing each subset in parallel

```
sum = 0;  
for (i = 0; i<1000; i = i + 1)  
    sum = sum + A[i];    /* sum local array subset
```

- The processors then coordinate in adding together the sub sums (P_n is the number of processors, `send(x,y)` sends value `y` to processor `x`, and `receive()` receives a value)


```
half = 100;  
limit = 100;  
repeat  
    half = (half+1)/2;    /*dividing line  
    if (Pn>= half && Pn<limit) send(Pn-half,sum);  
    if (Pn<(limit/2)) sum = sum + receive();  
    limit = half;  
until (half == 1);    /*final sum in P0's sum
```

* Slide is a courtesy of Mary Jane Irwin



An Example with 10 Processors

(Message Passing Model)



Pros and Cons of Message Passing

- ❑ Message sending and receiving is *much* slower than addition, for example
- ❑ But message passing multiprocessors are much easier for hardware designers to design
 - ➔ Don't have to worry about cache coherency for example
- ❑ The advantage for programmers is that communication is explicit, so there are fewer “performance surprises” than with the implicit communication in cache-coherent SMPs.
 - ➔ Message passing standard MPI-2 (www.mpi-forum.org)
- ❑ However, it is harder to port a sequential program to a message passing multiprocessor since every communication must be identified in advance.
 - ➔ With cache-coherent shared memory the hardware figures out what data needs to be communicated



Fundamental Issues

① Naming: how to solve large problem fast

- ➔ what data is shared
- ➔ how it is addressed
- ➔ what operations can access data
- ➔ how processes refer to each other

- ❑ Choice of naming affects code produced by a compiler; via load where just remember address or keep track of processor number and local virtual address for message passing
- ❑ Choice of naming affects replication of data; via load in cache memory hierarchy or via SW replication and consistency
- ❑ Global physical address space: any processor can generate, address and access it in a single operation
- ❑ Global virtual address space: if the address space of each process can be configured to contain all shared data of the parallel program
 - ➔ memory can be anywhere: virtual address translation handles it
- ❑ Segmented shared address space: locations are named <process number, address>
 - ➔ uniformly for all processes of the parallel program



Fundamental Issues

② Synchronization: To cooperate, processes must coordinate

- ❑ Message passing is implicit coordination with transmission or arrival of data
- ❑ Shared address => additional operations to explicitly coordinate:
e.g., write a flag, awaken a thread, interrupt a processor

③ Latency and Bandwidth

❑ Bandwidth

- ➔ Need high bandwidth in communication
- ➔ Cannot scale, but stay close
- ➔ Match limits in network, memory, and processor
- ➔ Overhead to communicate is a problem in many machines

❑ Latency

- ➔ Affects performance, since processor may have to wait
- ➔ Affects ease of programming, since requires more thought to overlap communication and computation

❑ Latency Hiding

- ➔ How can a mechanism help hide latency?
- ➔ Examples: overlap message send with computation, pre-fetch data, switch to other tasks

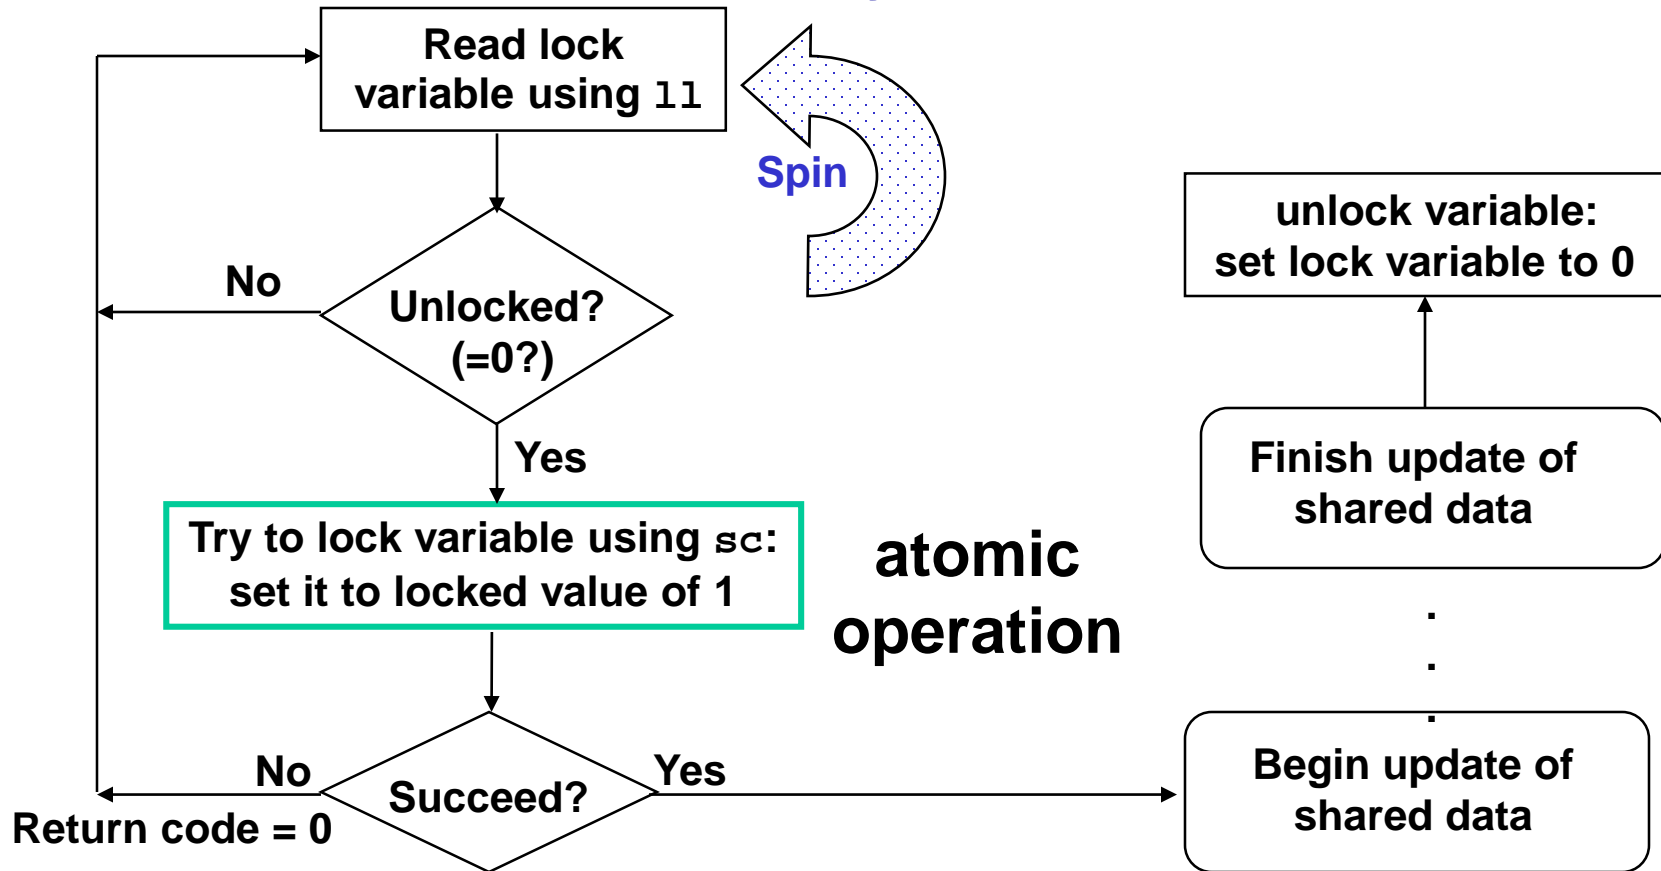


* Slide is a courtesy of Dave Patterson

Synchronization

- ❑ Why synchronize?
 - ➔ Need to know when it is safe for different processes to use shared data
 - ➔ To preserve execution semantics of multiprocessor applications
- ❑ Synchronization mechanisms are typically built with user-level software routines (e.g. *locks* and *barriers*) that rely on hardware-supplied primitives
- ❑ The key hardware-supplied synchronization capability is an uninterruptible instruction to atomically retrieve and change a value
- ❑ The implementation of hardware synchronization instructions introduces additional level of design complexity
- ❑ Single bus provides arbitration mechanism, since the bus is the only path to memory – the processor that gets the bus wins
- ❑ In large-scale machines or high-contention situations, synchronization can be a performance bottleneck due to the latency introduced by contention
- ❑ Issues for synchronization:
 - ➔ Negative impact on the execution time of the multiprocessor application
 - ➔ Interactions with cache coherence protocols (avoiding deadlocks)
 - ➔ Interface with user level programs

Spin Lock Synchronization



- ❑ Locking can be done via an **atomic swap operation**. MIPS has load-link (ll) and store-conditional (sc) instructions. Using both instructions a processor reads a location *and* sets it to locked state – **test-and-set** – in the same bus operation.
- ❑ The *single winning* processor will succeed in writing a 1 to the lock variable - all others processors will get a return code of 0

The Problem of Cache Coherency

- ❑ A cache coherence problem arises when the cache reflects a view of memory which is different from reality
- ❑ Cache coherence is a common issue when handling the I/O subsystem
- ❑ For the centralized shared memory architecture two different processors can view two different values for the same memory location

Time	Event	Cache Contents for CPU A	Cache Contents for CPU B	Memory Contents for location X
0				1
1	CPU A reads X	1		1
2	CPU B reads X	1	1	1
3	CPU A stores 0 into X	0	1	0

❑ A memory system is coherent if:

- ➊ A read by a processor P, to a location X follows a write by P to X, with no writes of X by another processor occurring in between, always returns the value written by P
- ➋ A read by P to location X that follows a write by another processor to X returns the newly written value if the read and write are sufficiently separated
- ➌ Writes to the same location are serialized: that is two writes to the same location by any two processors are seen in the same order by all processors (consistency)

HW Coherency Solutions

- ❑ In a coherent multiprocessor the cache provides both migration and replication of shared data
- ❑ Data migration allows low latency access to shared data from a local cache
- ❑ Replicating data enables simultaneous access to shared data and limits the potential of contention

Snooping Solution (Snoopy Bus)

- ➔ Send all requests for data to all processors
- ➔ Processors snoop to see if they have a copy and respond accordingly
- ➔ Requires broadcast, since caching information is at processors
- ➔ Works well with bus (natural broadcast medium)
- ➔ Dominates for small scale machines (most of the market)

Directory-Based Schemes

- ➔ Keep track of what is being shared in one centralized place
- ➔ Distributed memory => distributed directory for scalability (avoids bottlenecks)
- ➔ Send point-to-point requests to processors via network
- ➔ Scales better than Snooping
- ➔ Actually existed before Snooping-based schemes

* Slide is partially a courtesy of Dave Patterson



Conclusion

□ Summary

- ➔ Taxonomy of parallel architecture
 - Flynn's categorization (SISD, SIMD, MISD, MIMD)
 - Level of parallelism (bit level, instruction level, thread level)
 - MIMD architectures (centralized shared memory, distributed memory)
 - Programming and communication models
- ➔ Opportunities for applications
 - Scientific applications (weather, pharmaceuticals, aerospace, etc.)
 - Commercial applications (transaction processing and web servers)
- ➔ Communication models
 - Distributed shared memory model
 - Message passing model
- ➔ Fundamental issues for parallel computing
 - Naming of nodes and data items
 - Synchronization of application execution on the different nodes
 - Managing latency and bandwidth

□ Next Lecture

- ➔ Review

Read sections 7.1-7.6 in of the 4th Ed. Of textbook