

# CMPE 411

## Computer Architecture

### Lecture 20

## Memory Hierarchy & Cache

November 7, 2017

[www.csee.umbc.edu/~younis/CMPE411/  
CMPE411.htm](http://www.csee.umbc.edu/~younis/CMPE411/CMPE411.htm)



# Lecture's Overview

## Previous Lecture:

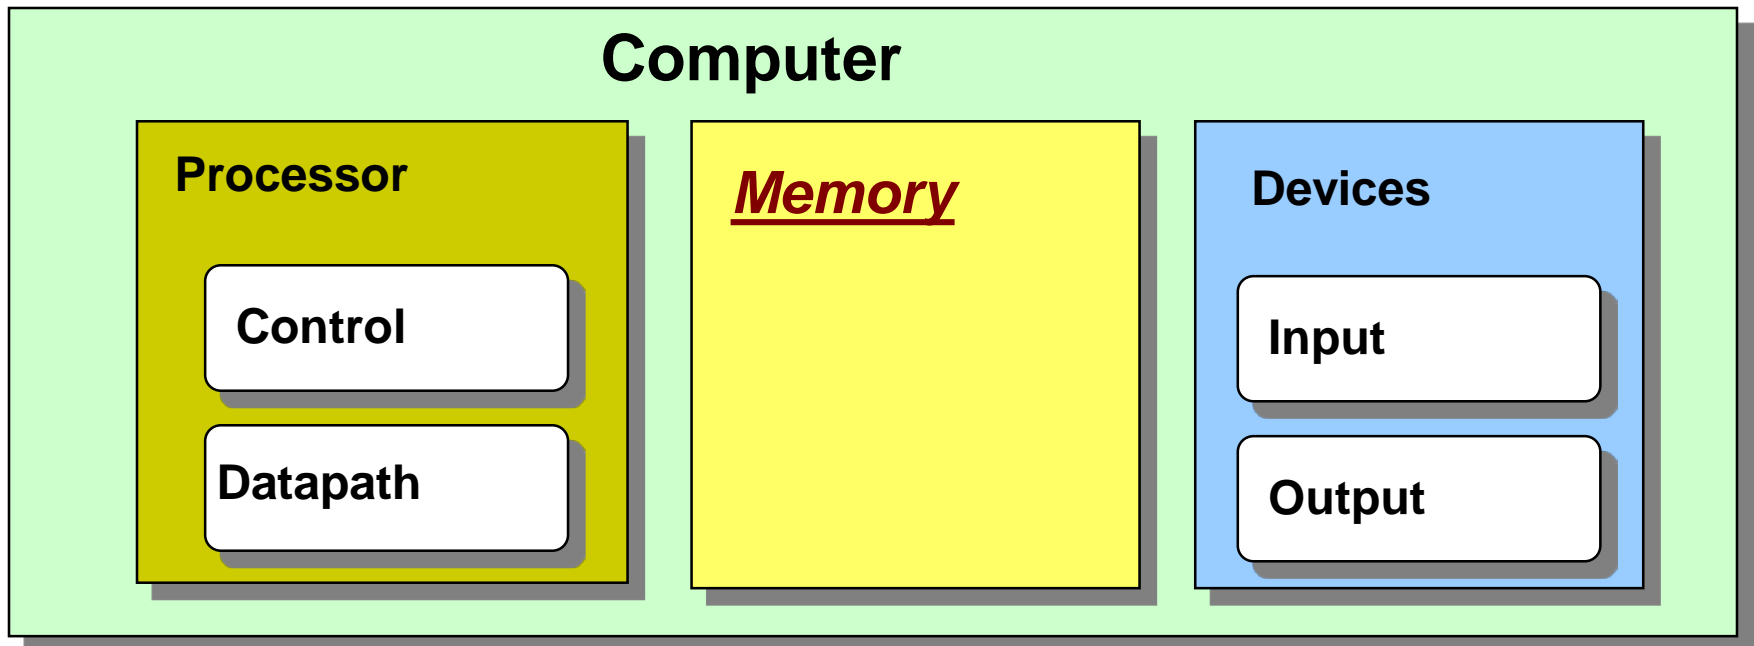
- Pipeline hazard detection
  - ➔ Data hazard
  - ➔ Control hazard
- Controlling pipeline operations
  - ➔ Data forwarding and branch prediction
  - ➔ Coordination between control and data hazard
- Supporting exceptions
  - ➔ Pipeline flushing
  - ➔ Invoking exception handling routines

## This Lecture:


- Memory hierarchy
- The basic of cache memory
- Organization of main memory

# Introduction

- ❑ Why do designers need to know about Memory technology?
  - ➔ Processor performance is usually limited by memory bandwidth
  - ➔ As IC densities increase, lots of memory will fit on processor chip
- ❑ What are the different types of memory?
- ❑ Why are not memory just a bunch of flip-flops?



# Processor-Memory Performance



## Problem:

Memory can be a bottleneck for processor performance

## Solution:

Rely on memory hierarchy of faster memory to bridge the gap

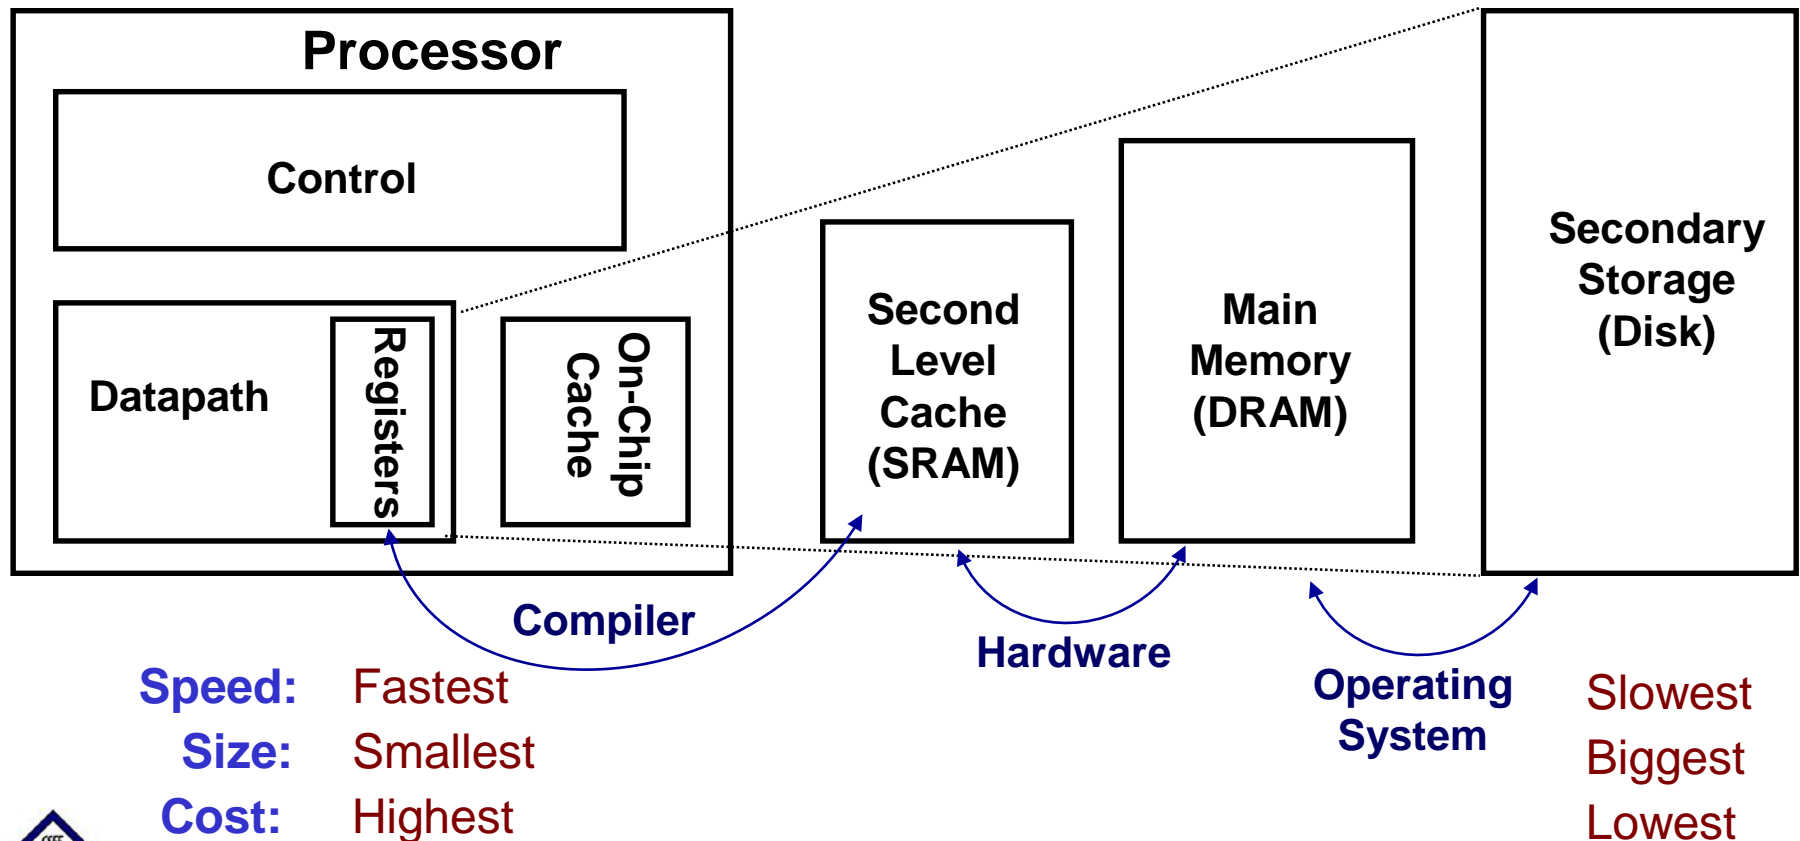
# Memory Hierarchy

## □ Temporal Locality (Locality in Time):

⇒ Keep most recently accessed data items closer to the processor


## □ Spatial Locality (Locality in Space):

⇒ Move blocks of contiguous words to the faster levels



# Memory Hierarchy Terminology

- ❑ **Hit:** data appear in some block in the faster level (example: Block X)
  - ➔ **Hit Rate:** the fraction of memory access found in the faster level
  - ➔ **Hit Time:** Time to access the faster level which consists of  
Memory access time + Time to determine hit/miss
- ❑ **Miss:** data need to be retrieve from a block in the slower level (Block Y)
  - ➔ **Miss Rate** =  $1 - (\text{Hit Rate})$
  - ➔ **Miss Penalty:** Time to replace a block in the upper level +  
Time to deliver the block the processor
- ❑ **Hit Time  $\ll$  Miss Penalty**



# Memory Access Types

## □ Random Access:

- “Random” is good: access time is the same for all locations
- **DRAM:** Dynamic Random Access Memory
  - High density, low power, cheap, slow
  - Dynamic: need to be “refreshed” regularly
- **SRAM:** Static Random Access Memory
  - Low density, high power, expensive, fast
  - Static: content will last “forever”(until losing power)

## □ “Non-so-random” Access:

- Access time varies from location to location and from time to time
- Examples: Disk, CD-ROM

## □ Sequential Access:

- Access time linear in location (e.g., Tape)
- Very slow, inflexible, cheap, commonly used for backups

## □ The focus of this and next lecture on random access technology

- The Main Memory: DRAMs + Caches: SRAMs



\* Slide is courtesy of Dave Patterson

# The Basics of Cache

- ❑ Cache is the name chosen to represent the first level of hierarchy between the CPU and main memory in the first commercial machine
- ❑ Caches first appeared in research machines in early 1960s
- ❑ Virtually every general-purpose computer produced today includes cache

Requesting  $X_n$  generates a miss and the word  $X_n$  will be brought from main memory to cache

X4
X1
$X_n - 2$
$X_n - 1$
X2
X3

a. Before the reference to  $X_n$

X4
X1
$X_n - 2$
$X_n - 1$
X2
$X_n$
X3

b. After the reference to  $X_n$


## Issues:

- ➔ How do we know that a data item is in cache?
- ➔ If so, How to find it?



# Direct-Mapped Cache

- ❑ A tag per cache block identifies the memory address it is mapped to
- ❑ A valid bit per cache block indicates whether the content is current and active



$$\text{Cache block address} = (\text{Block address}) \bmod (\text{Number of cache blocks})$$

# Accessing Cache

Cache Size depends on:

- # cache blocks
- # address bits
- Word size

Example:

For  $n$ -bit address, 4-byte word & 1024 cache blocks:

cache size =

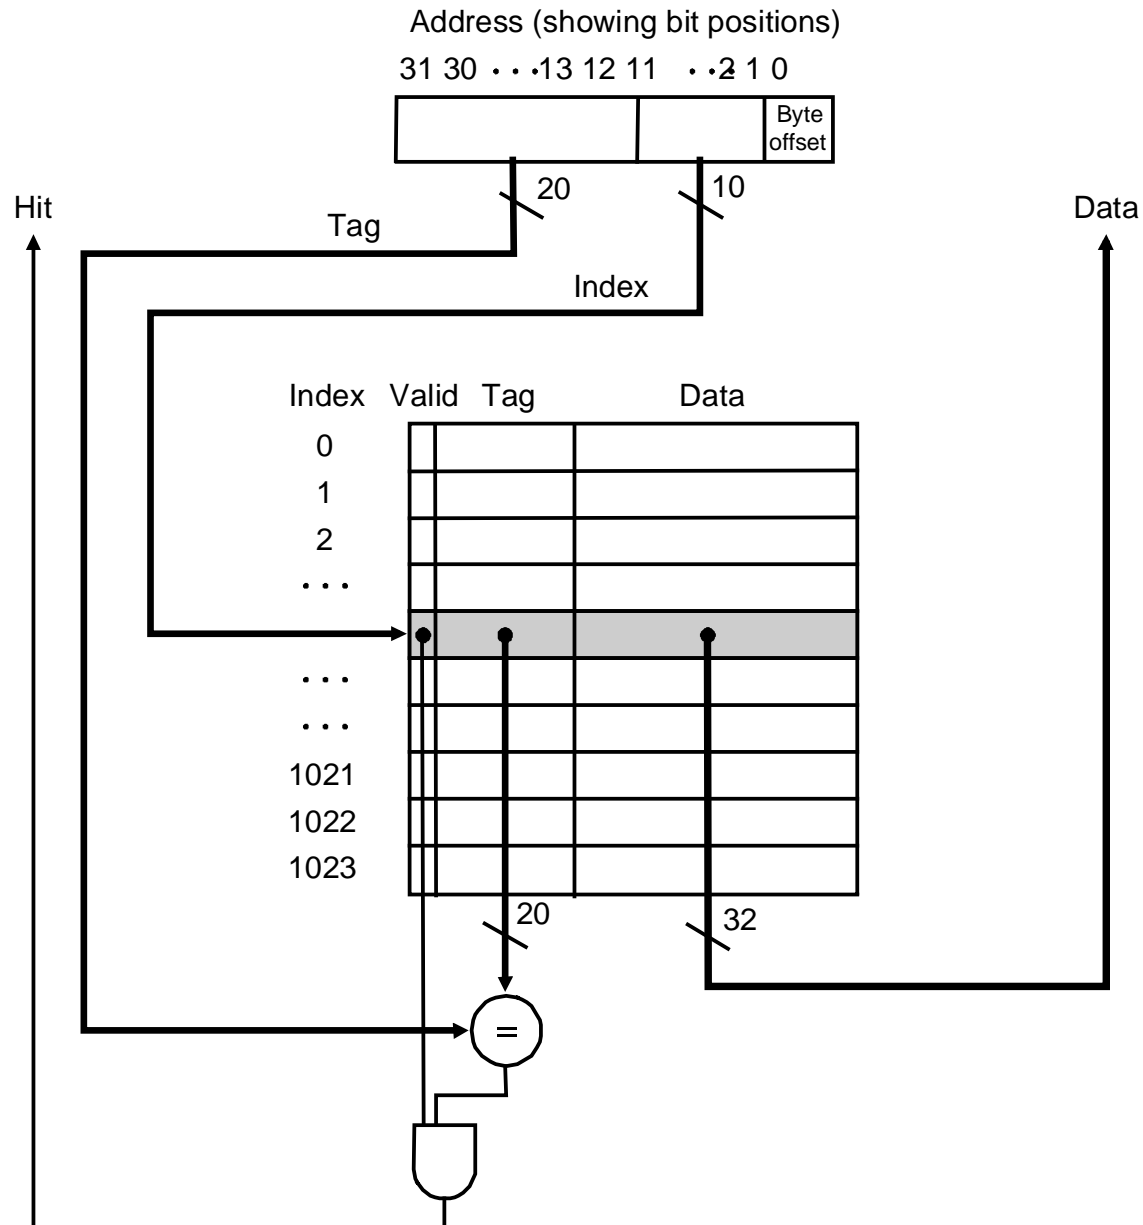
$$1024 [(n-10-2) + 1 + 32] \text{ bit}$$

Valid bit

Word size

Tag

# cache blocks



# Example

Reference address (Decimal)	Reference address (Binary)	Hit / miss	Assigned cache block
22	10110	Miss	$(10110 \bmod 1000) = 110$
26	11010	Miss	$(11010 \bmod 1000) = 010$
22	10110	Hit	$(10110 \bmod 1000) = 110$
26	11010	Hit	$(11010 \bmod 1000) = 010$

Index	Valid	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10	Memory(10110)
111	N		

Index	Valid	Tag	Data
000	N		
001	N		
010	Y	11	Memory(11010)
011	N		
100	N		
101	N		
110	Y	10	Memory(10110)
111	N		


# Example

Reference address (Decimal)	Reference address (Binary)	Hit / miss	Assigned cache block
22	10110	Miss	$(10110 \bmod 1000) = 110$
26	11010	Miss	$(11010 \bmod 1000) = 010$
22	10110	Hit	$(10110 \bmod 1000) = 110$
26	11010	Hit	$(11010 \bmod 1000) = 010$
16	10000	Miss	$(10000 \bmod 1000) = 000$
3	00011	Miss	$(00011 \bmod 1000) = 011$
16	10000	Hit	$(10000 \bmod 1000) = 000$

Index	Valid	Tag	Data
000	Y	10	Memory(10000)
001	N		
010	Y	11	Memory(11010)
011	N		
100	N		
101	N		
110	Y	10	Memory(10110)
111	N		

Index	Valid	Tag	Data
000	Y	10	Memory(10000)
001	N		
010	Y	11	Memory(11010)
011	Y	00	Memory(00011)
100	N		
101	N		
110	Y	10	Memory(10110)
111	N		

# Cache with Multi-Word/Block



- Takes advantage of spatial locality to improve performance
- *Cache block address = (Block address) modulo (Number of cache blocks)*
- *Block address = (byte address)/(bytes per block)*

# Determining Block Size

❑ Larger block size take advantage of spatial locality **BUT**:




➔ Larger block size means larger miss penalty:

- Takes longer time to fill up the block

➔ If block size is too big relative to cache size, miss rate will go up

- Too few cache blocks

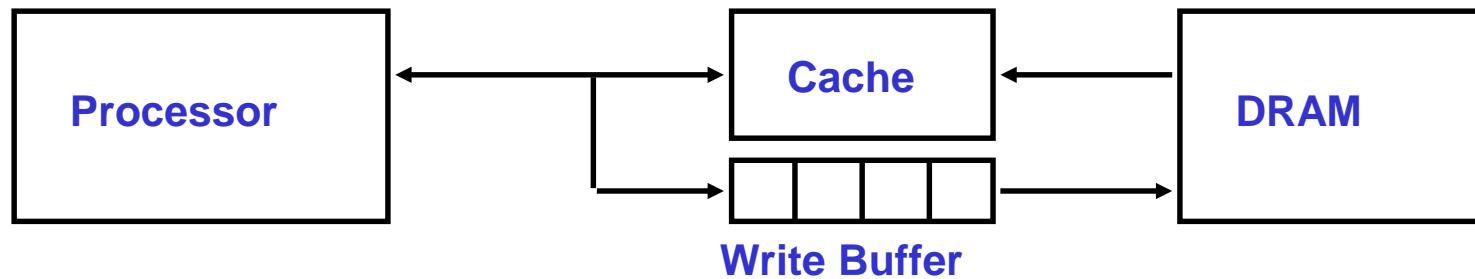
$$\text{Average Access Time} = \text{Hit Time} \times (1 - \text{Miss Rate}) + \text{Miss Penalty} \times \text{Miss Rate}$$




# Handling Cache Misses

- ❑ Misses for read access always bring blocks from main memory to cache
- ❑ Write access requires careful maintenance of consistency between cache and main memory
- ❑ Two possible strategies for handling write access misses:
  - ① Write through: The information is written to both the block in the cache and to the block in the slower memory
    - ➔ Read misses cannot result in writes
    - ➔ No allocation of a cache block is needed
    - ➔ Always combined with write buffers so that don't wait for slow memory
  - ② Write back: The information is written only to the block in the cache. The modified cache block is written to main memory only when it is replaced
    - ➔ Is block clean or dirty?
    - ➔ No writes to slow memory for repeated write accesses
    - ➔ Requires allocation of a cache block

# Write Through via Buffering




- ☐ Processor writes data into the cache and the write buffer
- ☐ Memory controller writes contents of the buffer to memory
- ☐ Increased write frequency can cause saturation of write buffer
- ☐ If CPU cycle time too fast and/or too many store instr. in a row:
  - ➔ Store buffer will overflow no matter how big you make it
  - ➔ The CPU Cycle Time get closer to DRAM Write Cycle Time
- ☐ Write buffer saturation can be handled by installing a second level (L2) cache







# Memory Organization



a. One-word-wide memory organization



b. Wide memory organization



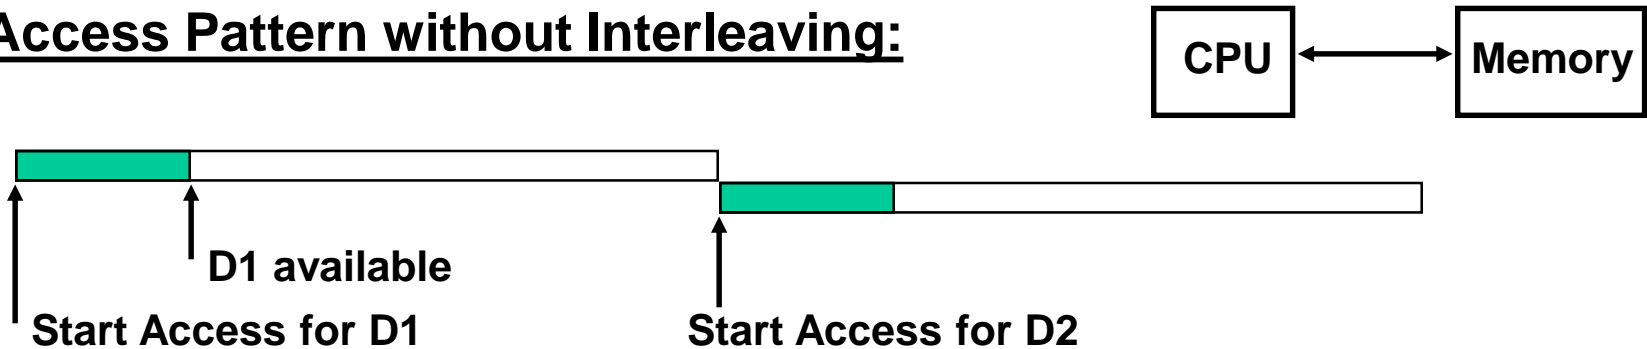
c. Interleaved memory organization

- ❑ **Simple:** CPU, Cache, Bus, Memory same width (32 bits)
- ❑ **Wide:** CPU/Mux 1 word; Mux/Cache, Bus, Memory N words
- ❑ **Interleaved:** CPU, Cache, Bus 1 word: Memory N Modules (4 Modules); example is *word interleaved*


**Memory organization would have significant effect on bandwidth**

# Memory Interleaving

## Access Pattern without Interleaving:



## Access Pattern with 4-way Interleaving:



# Conclusion

## □ Summary

- ➔ Memory hierarchy
  - Principles of locality
  - Types of memory
- ➔ The basics of cache memory
  - Direct-mapped cache
  - Handling of cache misses
  - Consistency between cache and main memory
- ➔ Memory organization
  - Main memory performance issues
  - Memory interleaving

## □ Next Lecture

- ➔ Measuring and improving cache performance

**Read sections 5.1 & 5.2 in 4<sup>th</sup> Ed. Or 5<sup>th</sup> Ed. of the textbook**