

# Sabbir Ahmed

CMSC 421: Project 2 Final Design Document

May 13, 2018

---

## 1 Introduction

This project implements a new version of the Linux kernel that adds functionality to support a simple intrusion detection system (IDS). This system will operate by logging the system calls made by a process in the kernel, while analysis and intrusion detection will be done in user space. This assignment is designed to teach a simple method of intrusion detection, as well as to reinforce the idea of how user space and kernel space interact through the use of system calls.

An intrusion detection system is a computer program that attempts to identify (and thwart) attacks that might be performed on the system by attackers. There are several time-tested approaches to the development of an IDS. The project will keep track of the system calls made by a monitored process and check for abnormalities in the sequences of system calls made. When an attacker breaks into a process, they will need to make system calls in order to attempt to access the resources of the system that are under attack. As the system calls that the attacker will perform will likely be different than those performed by a process that is not under attack, it follows that by monitoring both healthy and broken processes, it is possible to develop a scheme to identify those that might be under attack for further action to be taken.

## 2 Objective

The project will compare sequences of system calls made by a monitored process to known good sequences.

### 2.1 Kernel Space Requirements

The kernel-space program will instrument the system call dispatcher of the Linux kernel with code that logs each time a system call is made. Built-in system calls such as `ptrace` are not allowed to trace the usage of system calls

to generate the logs for the project.

## **2.2 User Space Requirements**

The analysis of the logs will be handled by the user-space program, which may be implemented in any supported programming language. The user-space process should construct a bit array for each process under monitoring showing which system calls have been run in a window of the last  $k$  system calls. If a particular system call is made in the window, the bit for that system call will be set to 1. The bit arrays will then be measured for their hamming distance with the example of a healthy system call sequence for a process.

## **3 Design Approach**

The final design of the project adopted a much simpler approach than what was previously assumed. The methods to achieve the requirements became apparent after extensive research.

### **3.1 Kernel Space**

The kernel-space program will trace the system calls and dump them on logs consisting of the pid and the system call numbers of each processes spawned by a program. The log was intended to also contain the corresponding timestamps of the execution of the system calls, but the kernel-space did not appear to have nanosecond resolution of their timestamps. Initially, the kernel-space program intended to create the log files in  $N$  chunks specified by the user to pass on to the user-space program for analysis. This parameter will be removed as an option to the user and instead be converted to a constant. The constant  $N$  number of system calls per log files would provide a smooth communication method with the user-space program.

Initially, `ptrace` appeared to be the only method of tracing system calls. Two different approaches were proposed in the initial design. The first approach suggested inserting breakpoints in the implementations for each system calls in the kernel. The second and more probable approach suggested a stripped down version of `ptrace` that behaved identically. After extensive research, a third approach was considered that merged aspects of both of the

prior approaches. The conditional jump instruction, `jnz`, was updated to an unconditional jump, `jmp`, inside the system call entry point section in `entry_64.S`. This update forces all the system calls to save their system call number to registers before disabling the interrupt request. The system call number is then accessible in the register `ORIG_RAX` in `common.c` which can be provided as a macro to external files.

The kernel-space program will require the user to provide the programs to be traced. The program will fork its own copy of the process, and gather the process IDs and the system call numbers. The `proc` filesystem will be utilized to provide a convenient interface between the kernel-space program and writing log files accessible to the user space. The common directory shared by the entire IDS, `idsdata`, will be used to store the log files.

## **3.2 User Space**

The user-space program was developed before the kernel-space program. The program was tested using dummy data simulating log files created by the kernel-space program. An object oriented approach was taken to develop the interface with the user and the kernel-space logger. The user would be required to initialize the user-space program after the kernel-space program has been initiated to run concurrently. The user-space program would wait for the logger to generate a chunk of the system call logs, and then automatically store the first